

# Task-level quality

Mean overall score over 10 repeats per task; whiskers are 95% confidence intervals.

## Hard prompts separate robust wins from failure modes

